

Logistic_model

Chloe Florence and Tiffany Le

2025-09-26

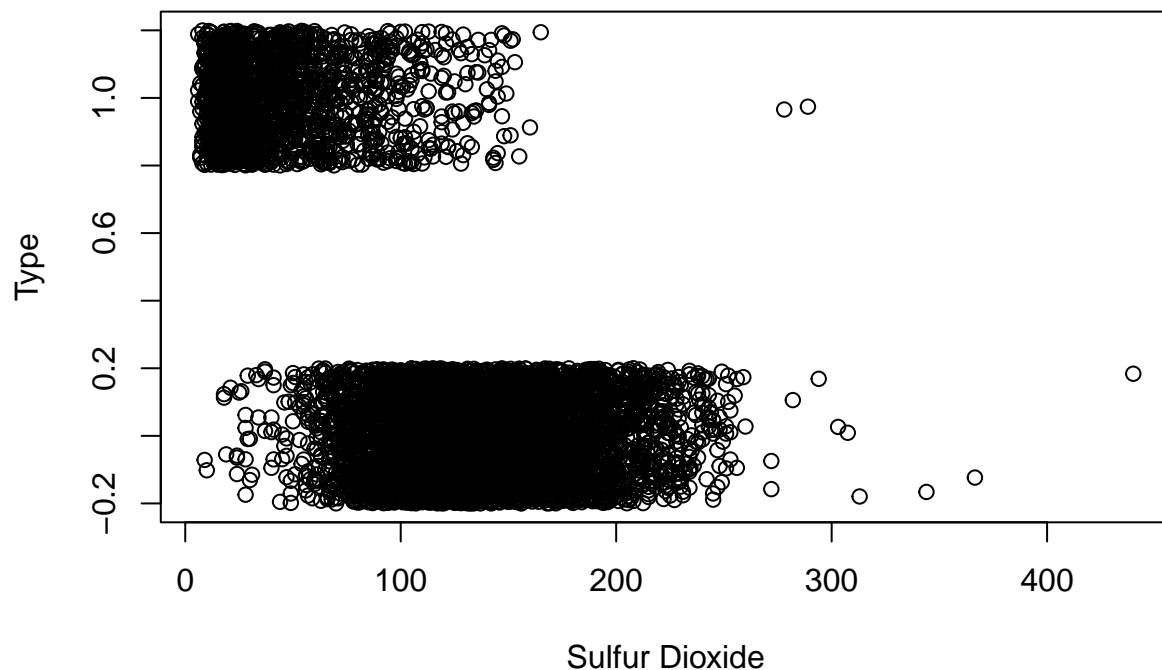
Importing Data

```
library(dplyr)
library(class)
library(leaps)
# creating binomial predictor for type of wine
red_wine <- read.csv("winequality-red.csv", header = TRUE, sep = ";")
red_wine["type"] = 1
white_wine <- read.csv("winequality-white.csv", header = TRUE, sep = ";")
white_wine["type"] = 0
# creating binomial predictor for if a wine is good or not.
wine <- bind_rows(red_wine, white_wine)
wine$good <- ifelse(wine$quality > 5, 1, 0)
```

Data Exploration

Plots

```
plot(wine$total.sulfur.dioxide, jitter(wine$type), xlab = "Sulfur Dioxide", ylab = "Type")
```




```
# looking at correlation of various predictors
cor(wine)
```

```
##          fixed.acidity volatile.acidity citric.acid residual.sugar
## fixed.acidity      1.00000000      0.21900826  0.32443573   -0.11198128
## volatile.acidity    0.21900826      1.00000000 -0.37798132   -0.19601117
## citric.acid         0.32443573     -0.37798132  1.00000000    0.14245123
## residual.sugar     -0.11198128     -0.19601117  0.14245123    1.00000000
## chlorides          0.29819477      0.37712428  0.03899801   -0.12894050
## free.sulfur.dioxide -0.28273543     -0.35255731  0.13312581    0.40287064
## total.sulfur.dioxide -0.32905390     -0.41447619  0.19524198    0.49548159
## density            0.45890998      0.27129565  0.09615393    0.55251695
## pH                 -0.25270047      0.26145440 -0.32980819   -0.26731984
## sulphates          0.29956774      0.22598368  0.05619730   -0.18592741
## alcohol            -0.09545152     -0.03764039 -0.01049349   -0.35941477
## quality            -0.07674321     -0.26569948  0.08553172   -0.03698048
## type               0.48673983      0.65303559 -0.18739650   -0.34882101
## good              -0.06735375     -0.26704633  0.07573859   -0.03248435
##          chlorides free.sulfur.dioxide total.sulfur.dioxide
## fixed.acidity      0.29819477      -0.28273543      -0.32905390
## volatile.acidity    0.37712428      -0.35255731      -0.41447619
## citric.acid         0.03899801      0.13312581      0.19524198
## residual.sugar     -0.12894050      0.40287064      0.49548159
## chlorides          1.00000000      -0.19504479      -0.27963045
## free.sulfur.dioxide -0.19504479      1.00000000      0.72093408
## total.sulfur.dioxide -0.27963045      0.72093408      1.00000000
## density            0.36261466      0.02571684      0.03239451
## pH                 0.04470798      -0.14585390      -0.23841310
## sulphates          0.39559331      -0.18845725      -0.27572682
## alcohol            -0.25691558      -0.17983843      -0.26573964
## quality            -0.20066550      0.05546306      -0.04138545
## type               0.51267825      -0.47164366      -0.70035716
## good              -0.18190812      0.04481948      -0.04758506
##          density      pH      sulphates      alcohol
## fixed.acidity      0.45890998 -0.25270047  0.29956774 -0.095451523
## volatile.acidity    0.27129565  0.26145440  0.22598368 -0.037640386
## citric.acid         0.09615393 -0.32980819  0.05619730 -0.010493492
## residual.sugar     0.55251695 -0.26731984 -0.185927405 -0.359414771
## chlorides          0.36261466  0.04470798  0.395593307 -0.256915580
## free.sulfur.dioxide 0.02571684 -0.14585390 -0.188457249 -0.179838435
## total.sulfur.dioxide 0.03239451 -0.23841310 -0.275726820 -0.265739639
## density            1.00000000  0.01168608  0.259478495 -0.686745422
## pH                 0.01168608  1.00000000  0.192123407  0.121248467
## sulphates          0.25947850  0.19212341  1.000000000 -0.003029195
## alcohol            -0.68674542  0.12124847 -0.003029195  1.000000000
## quality            -0.30585791  0.01950570  0.038485446  0.444318520
## type               0.39064532  0.32912865  0.487217970 -0.032969551
## good              -0.26887620  0.01884228  0.035807167  0.394675617
##          quality      type      good
## fixed.acidity      -0.07674321  0.48673983 -0.06735375
## volatile.acidity    -0.26569948  0.65303559 -0.26704633
## citric.acid         0.08553172 -0.18739650  0.07573859
## residual.sugar     -0.03698048 -0.34882101 -0.03248435
## chlorides          -0.20066550  0.51267825 -0.18190812
```

```
## free.sulfur.dioxide  0.05546306 -0.47164366  0.04481948
## total.sulfur.dioxide -0.04138545 -0.70035716 -0.04758506
## density             -0.30585791  0.39064532 -0.26887620
## pH                  0.01950570  0.32912865  0.01884228
## sulphates           0.03848545  0.48721797  0.03580717
## alcohol             0.44431852 -0.03296955  0.39467562
## quality             1.00000000 -0.11932328  0.81448366
## type                -0.11932328  1.00000000 -0.11659486
## good                0.81448366 -0.11659486  1.00000000
```

```
boxplot(wine$total.sulfur.dioxide, wine$type, xlab = "Type of Wine", ylab = "Total Sulfur Dioxide")
```




```
# looking for possible outliers
```

```
wine[wine[, "total.sulfur.dioxide"] > 400,]
```

```
##      fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 6345          6.1           0.26         0.25           2.9         0.047
##      free.sulfur.dioxide total.sulfur.dioxide density   pH sulphates alcohol
## 6345                289                440 0.99314 3.44         0.64     10.5
##      quality type good
## 6345        3    0    0
```


```
plot(wine$chlorides, jitter(wine$type), xlab = "Chlorides", ylab = "Type")
```




```
# possible outliers
wine[wine[, "chlorides"] > 0.6,]
```

```
##      fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 152          9.2           0.52         1.00          3.4      0.610
## 259          7.7           0.41         0.76          1.8      0.611
##      free.sulfur.dioxide total.sulfur.dioxide density    pH sulphates alcohol
## 152                32                69 0.9996 2.74      2.00      9.4
## 259                8                45 0.9968 3.06      1.26      9.4
##      quality type good
## 152         4     1    0
## 259         5     1    0
```

```
# evaluating effect of outliers
plot(wine$alcohol, wine$density, xlab = "Alcohol", ylab = "Density")
```



```
wo_outlier <- wine[wine["density"] < 1.01,]
plot(wo_outlier$alcohol, wo_outlier$density, xlab = "Alcohol", ylab = "Density")
```



```
wine[wine["density"] > 1.01,]
```

```
##      fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 3253          7.9          0.330         0.28          31.6       0.053
## 3263          7.9          0.330         0.28          31.6       0.053
## 4381          7.8          0.965         0.60          65.8       0.074
##      free.sulfur.dioxide total.sulfur.dioxide density    pH sulphates alcohol
## 3253                35                176 1.01030 3.15      0.38      8.8
```

```
## 3263          35          176 1.01030 3.15      0.38      8.8
## 4381          8          160 1.03898 3.39      0.69     11.7
##      quality type good
## 3253      6      0      1
## 3263      6      0      1
## 4381      6      0      1
```

Summary Statistics

```
summary(wine)
```

```
## fixed.acidity    volatile.acidity    citric.acid      residual.sugar
## Min.   : 3.800    Min.   :0.0800    Min.   :0.0000    Min.   : 0.600
## 1st Qu.: 6.400    1st Qu.:0.2300    1st Qu.:0.2500    1st Qu.: 1.800
## Median : 7.000    Median :0.2900    Median :0.3100    Median : 3.000
## Mean   : 7.215    Mean   :0.3397    Mean   :0.3186    Mean   : 5.443
## 3rd Qu.: 7.700    3rd Qu.:0.4000    3rd Qu.:0.3900    3rd Qu.: 8.100
## Max.   :15.900    Max.   :1.5800    Max.   :1.6600    Max.   :65.800
## chlorides      free.sulfur.dioxide    total.sulfur.dioxide    density
## Min.   :0.00900    Min.   : 1.00      Min.   : 6.0          Min.   :0.9871
## 1st Qu.:0.03800    1st Qu.: 17.00      1st Qu.: 77.0         1st Qu.:0.9923
## Median :0.04700    Median : 29.00      Median :118.0         Median :0.9949
## Mean   :0.05603    Mean   : 30.53      Mean   :115.7         Mean   :0.9947
## 3rd Qu.:0.06500    3rd Qu.: 41.00      3rd Qu.:156.0         3rd Qu.:0.9970
## Max.   :0.61100    Max.   :289.00      Max.   :440.0         Max.   :1.0390
##      pH          sulphates          alcohol          quality
## Min.   :2.720    Min.   :0.2200    Min.   : 8.00      Min.   :3.000
## 1st Qu.:3.110    1st Qu.:0.4300    1st Qu.: 9.50      1st Qu.:5.000
## Median :3.210    Median :0.5100    Median :10.30      Median :6.000
## Mean   :3.219    Mean   :0.5313    Mean   :10.49      Mean   :5.818
## 3rd Qu.:3.320    3rd Qu.:0.6000    3rd Qu.:11.30      3rd Qu.:6.000
## Max.   :4.010    Max.   :2.0000    Max.   :14.90      Max.   :9.000
##      type          good
## Min.   :0.0000    Min.   :0.0000
## 1st Qu.:0.0000    1st Qu.:0.0000
## Median :0.0000    Median :1.0000
## Mean   :0.2461    Mean   :0.6331
## 3rd Qu.:0.0000    3rd Qu.:1.0000
## Max.   :1.0000    Max.   :1.0000
```

Logistic Model

```
set.seed(1)
train <- sample(6497, 5198)
# Logistic full model
log_mod <- glm(good ~ . - quality, data = wine, family = binomial, subset = train)
summary(log_mod)
```

```
##
## Call:
## glm(formula = good ~ . - quality, family = binomial, data = wine,
##      subset = train)
##
```

```
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    1.443e+02  4.907e+01   2.941  0.00327 **
## fixed.acidity    1.193e-01  5.587e-02   2.136  0.03272 *
## volatile.acidity -4.618e+00  3.263e-01 -14.152 < 2e-16 ***
## citric.acid     -3.759e-01  2.863e-01  -1.313  0.18920
## residual.sugar   1.284e-01  2.106e-02   6.098  1.07e-09 ***
## chlorides       -1.199e+00  1.146e+00  -1.047  0.29527
## free.sulfur.dioxide 1.201e-02  2.818e-03   4.262  2.03e-05 ***
## total.sulfur.dioxide -5.621e-03  1.172e-03  -4.795  1.62e-06 ***
## density         -1.561e+02  4.990e+01  -3.128  0.00176 **
## pH              8.970e-01  3.307e-01   2.712  0.00669 **
## sulphates        2.058e+00  2.956e-01   6.961  3.38e-12 ***
## alcohol         7.705e-01  6.638e-02  11.607 < 2e-16 ***
## type            5.750e-01  2.097e-01   2.742  0.00611 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 6825.3  on 5197  degrees of freedom
## Residual deviance: 5371.1  on 5185  degrees of freedom
## AIC: 5397.1
##
## Number of Fisher Scoring iterations: 4
# best predictors for each number of predictors

regfit.fwd = regsubsets(good ~ . - quality, data = wine, subset = train)
summary(regfit.fwd)

## Subset selection object
## Call: regsubsets.formula(good ~ . - quality, data = wine, subset = train)
## 12 Variables (and intercept)
##              Forced in Forced out
## fixed.acidity      FALSE      FALSE
## volatile.acidity    FALSE      FALSE
## citric.acid         FALSE      FALSE
## residual.sugar      FALSE      FALSE
## chlorides           FALSE      FALSE
## free.sulfur.dioxide FALSE      FALSE
## total.sulfur.dioxide FALSE      FALSE
## density             FALSE      FALSE
## pH                  FALSE      FALSE
## sulphates           FALSE      FALSE
## alcohol             FALSE      FALSE
## type                FALSE      FALSE
## 1 subsets of each size up to 8
## Selection Algorithm: exhaustive
##      fixed.acidity volatile.acidity citric.acid residual.sugar chlorides
## 1  ( 1 ) " "      " "              " "              " "      " "
## 2  ( 1 ) " "      "*"              " "              " "      " "
## 3  ( 1 ) " "      "*"              " "              " "      " "
## 4  ( 1 ) " "      "*"              " "      "*"      " "
## 5  ( 1 ) " "      "*"              " "      "*"      " "
```

```
## 6 ( 1 ) " "      "*"      " "      "*"      " "
## 7 ( 1 ) " "      "*"      " "      "*"      " "
## 8 ( 1 ) " "      "*"      " "      "*"      "*"
##      free.sulfur.dioxide total.sulfur.dioxide density pH sulphates alcohol
## 1 ( 1 ) " "      " "      " "      " " " "      "*"
## 2 ( 1 ) " "      " "      " "      " " " "      "*"
## 3 ( 1 ) " "      " "      " "      " " "*"      "*"
## 4 ( 1 ) " "      " "      " "      " " "*"      "*"
## 5 ( 1 ) " "      "*"      " "      " " "*"      "*"
## 6 ( 1 ) "*"      "*"      " "      " " "*"      "*"
## 7 ( 1 ) "*"      "*"      " "      "*" "*"      "*"
## 8 ( 1 ) "*"      "*"      " "      "*" "*"      "*"
##      type
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "
## 6 ( 1 ) " "
## 7 ( 1 ) " "
## 8 ( 1 ) " "
```

```
# Best model through Best Subset Selection with BIC criteria
library(bestglm)
model <- bestglm(wine[train,-12], IC="BIC")
model
```

```
## BIC
## BICq equivalent for q in (5.45321160237977e-05, 0.635151358655417)
## Best Model:
##      Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.913098797 0.0718819304 -12.702758 1.980067e-36
## volatile.acidity -0.829760730 0.0395801022 -20.964088 9.725677e-94
## residual.sugar 0.011733475 0.0014560232 8.058576 9.499706e-16
## free.sulfur.dioxide 0.002507038 0.0004718694 5.312991 1.123437e-07
## total.sulfur.dioxide -0.001324350 0.0001655762 -7.998434 1.541554e-15
## sulphates 0.305566994 0.0407209526 7.503925 7.246010e-14
## alcohol 0.160061285 0.0053200272 30.086554 1.865151e-183
```

finding the test error for logistic regression

```
test<- wine[-train,]
glm.probs <- predict(log_mod, test , type = "response")

glm.pred <- rep(0, 1299)
glm.pred[glm.probs > .5] <- 1

table(glm.pred, test$good)

##
## glm.pred 0 1
##      0 278 120
##      1 206 695
```



```
mean(glm.pred != test$good)
```

```
## [1] 0.2509623
```

The test error in this case was 25.096%

Cross Validation with the logistic model

We will be training the data on 4/5ths of the data and then using the remaining 1/5 to test the data. The cross validation error will be the mean of the 5 different test errors we will get.

```
# getting the 5 different samples
set.seed(1)
a <- sample(1:6497, 1299)
b <- sample(setdiff(1:6497, a), 1299)
c <- sample(setdiff(1:6497, c(a, b)), 1299)
d <- sample(setdiff(1:6497, c(a, b, c)), 1300)
e <- setdiff(1:6497, c(a, b, c, d))

library(pROC)
par(mfrow=c(2,3))
samples <- list(a = a, b = b, c = c, d = d, e = e)
error <- c()
for (i in 1:5){
  # model for this training data set
  test <- samples[[i]]
  train <- setdiff(1:nrow(wine), test)
  log_mod <- glm(good ~ volatile.acidity + residual.sugar + free.sulfur.dioxide + total.sulfur.dioxide,
    data = wine[train,], type = "response")
  glm.probs <- predict(log_mod, wine[test,], type = "response")

  glm.pred <- ifelse(glm.probs > 0.5, 1, 0)

  actual <- wine$good[test]
  new_err <- mean(glm.pred != actual)

  error <- c(error, new_err)

  roc_obj <- roc(actual, glm.pred)
  plot(roc_obj, main = "ROC Curve", xlab = "False Positive Rate (1 - Specificity)",
    ylab = "True Positive Rate (Sensitivity)",
    print.auc(roc_obj))
}
```

```
## Area under the curve: 0.7271
```

```
## Area under the curve: 0.6968
```

```
## Area under the curve: 0.7015
```

```
## Area under the curve: 0.6851
```

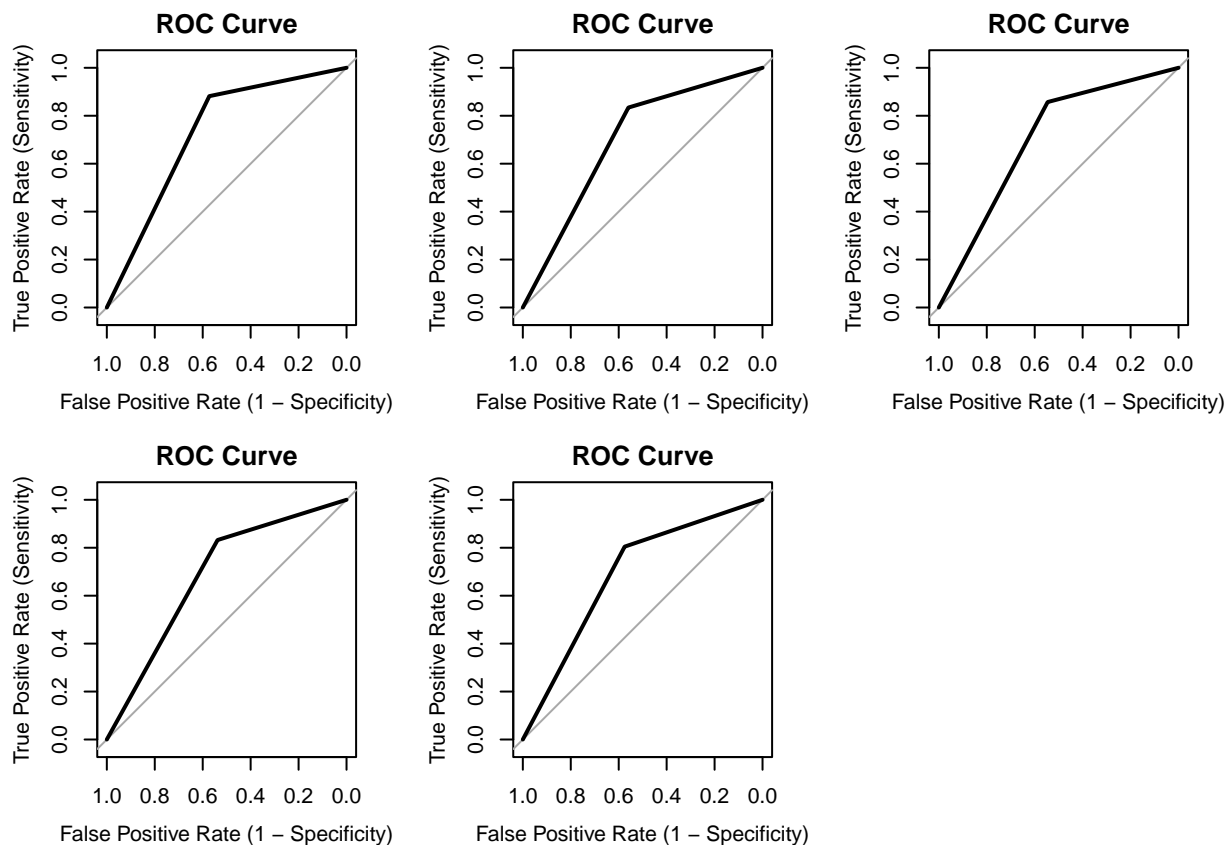
```
## Area under the curve: 0.6898
```

```
error
```

```
## [1] 0.2301771 0.2671286 0.2578907 0.2784615 0.2776923
```

```
mean(error)
```

```
## [1] 0.26227
```



accuracy = $1 - 0.26227 = 0.73773 \sim 73.77\%$ accurate

Precision Average = 0.766588708513

This means on average 76.65% of the true positive are correct out of the total predicted positive.

Recall: 0.841987326462

This means on average 84.12% of the true positive are correct out of the total true positive.

Full Logistic Model

```
full_mod <- log_mod <- glm(good ~ volatile.acidity + residual.sugar + free.sulfur.dioxide + total.sulfur.dioxide + sulphates + alcohol, family = binomial, data = wine)
summary(full_mod)
```

```
##
## Call:
## glm(formula = good ~ volatile.acidity + residual.sugar + free.sulfur.dioxide +
##     total.sulfur.dioxide + sulphates + alcohol, family = binomial,
##     data = wine)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -8.7440138   0.4201198  -20.813   < 2e-16 ***
## volatile.acidity  -4.3169708   0.2212239  -19.514   < 2e-16 ***
## residual.sugar     0.0633281   0.0074512    8.499   < 2e-16 ***
## free.sulfur.dioxide  0.0168408   0.0024883    6.768 1.31e-11 ***
## total.sulfur.dioxide -0.0076431   0.0008395   -9.105   < 2e-16 ***
## sulphates         1.8751919   0.2263165    8.286   < 2e-16 ***
```

```
## alcohol          0.9518199  0.0339175  28.063  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 8541.0  on 6496  degrees of freedom
## Residual deviance: 6723.1  on 6490  degrees of freedom
## AIC: 6737.1
##
## Number of Fisher Scoring iterations: 4
```

KNN Algorithm

```
# knn error with a as test set
train <- c(b, c, d, e)

knn_a <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual
table(knn_a, wine$good[a])

##
## knn_a    0    1
##      0 284 160
##      1 186 669

error_a <- mean(knn_a != wine$good[a]); error_a

## [1] 0.2663587

# knn error with b as test set
train <- c(a, c, d, e)

knn_b <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual
table(knn_b, wine$good[b])

##
## knn_b    0    1
##      0 289 161
##      1 190 659

error_b <- mean(knn_b != wine$good[b]); error_b

## [1] 0.2702079

# knn error with c as test set
train <- c(a, b, d, e)

knn_c <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual
table(knn_c, wine$good[c])

##
## knn_c    0    1
##      0 301 170
```

```

##      1 179 649
error_c <- mean(knn_c != wine$good[c]); error_c

## [1] 0.2686682
# knn error with d as test set
train <- c(a, b, c, e)

knn_d <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual
table(knn_d, wine$good[d])

##
## knn_d   0   1
##      0 321 164
##      1 168 647
error_d <- mean(knn_d != wine$good[d]); error_d

## [1] 0.2553846
# knn error with e as test set
train <- c(a, b, c, d)

knn_e <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual
table(knn_e, wine$good[e])

##
## knn_e   0   1
##      0 295 165
##      1 171 669
error_e <- mean(knn_e != wine$good[e]); error_e

## [1] 0.2584615
# mean of all 5 misclassification rates
mean(error_a, error_b, error_c, error_d, error_e)

## [1] 0.2663587
wine_knn_cv <- function(k) {
  train <- c(b, c, d, e)

  knn_a <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual
  error_a <- mean(knn_a != wine$good[a])

  train <- c(a, c, d, e)

  knn_b <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual
  error_b <- mean(knn_b != wine$good[b])

  train <- c(a, b, d, e)

  knn_c <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual

```

```

error_c <- mean(knn_c != wine$good[c])

train <- c(a, b, c, e)

knn_d <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual.sugar[train]), wine$volatile.acidity[-train], wine$residual.sugar[-train])

error_d <- mean(knn_d != wine$good[d])

train <- c(a, b, c, d)

knn_e <- knn(data.frame(volatile.acidity = wine$volatile.acidity[train], residual.sugar = wine$residual.sugar[train]), wine$volatile.acidity[-train], wine$residual.sugar[-train])

error_e <- mean(knn_e != wine$good[e])

mean(error_a, error_b, error_c, error_d, error_e)
}

```

KNN error rate for different values of k

```

# misclassification rate when k = 2
wine_knn_cv(2)

```

```
## [1] 0.3156274
```

```
wine_knn_cv(3)
```

```
## [1] 0.3071594
```

```
wine_knn_cv(4)
```

```
## [1] 0.3063895
```

```
wine_knn_cv(5)
```

```
## [1] 0.3010008
```

```
wine_knn_cv(6)
```

```
## [1] 0.3033102
```

```
wine_knn_cv(7)
```

```
## [1] 0.2994611
```

```
wine_knn_cv(8)
```

```
## [1] 0.3110085
```

```
wine_knn_cv(9)
```

```
## [1] 0.3125481
```

```
wine_knn_cv(10)
```

```
## [1] 0.321786
```

```

cv_error <- numeric(100)
for (i in 1:100) {
  cv_error[i] <- wine_knn_cv(i)
}

```

```

}
min(cv_error)

## [1] 0.2663587

This minimum error rate comes when k = 1

par(mfrow = c(2, 3))

plot(roc(wine$good[a], as.numeric(knn_a)), main = "ROC Curve", xlab = "False Positive Rate", ylab = "True Positive Rate")
plot(roc(wine$good[b], as.numeric(knn_b)), main = "ROC Curve", xlab = "False Positive Rate", ylab = "True Positive Rate")
plot(roc(wine$good[c], as.numeric(knn_c)), main = "ROC Curve", xlab = "False Positive Rate", ylab = "True Positive Rate")
plot(roc(wine$good[d], as.numeric(knn_d)), main = "ROC Curve", xlab = "False Positive Rate", ylab = "True Positive Rate")
plot(roc(wine$good[e], as.numeric(knn_e)), main = "ROC Curve", xlab = "False Positive Rate", ylab = "True Positive Rate")

#Finding the area under the curve for the ROC curve.
auc(roc(wine$good[a], as.numeric(knn_a)))

## Area under the curve: 0.7056

auc(roc(wine$good[b], as.numeric(knn_b)))

## Area under the curve: 0.7035

auc(roc(wine$good[c], as.numeric(knn_c)))

## Area under the curve: 0.7098

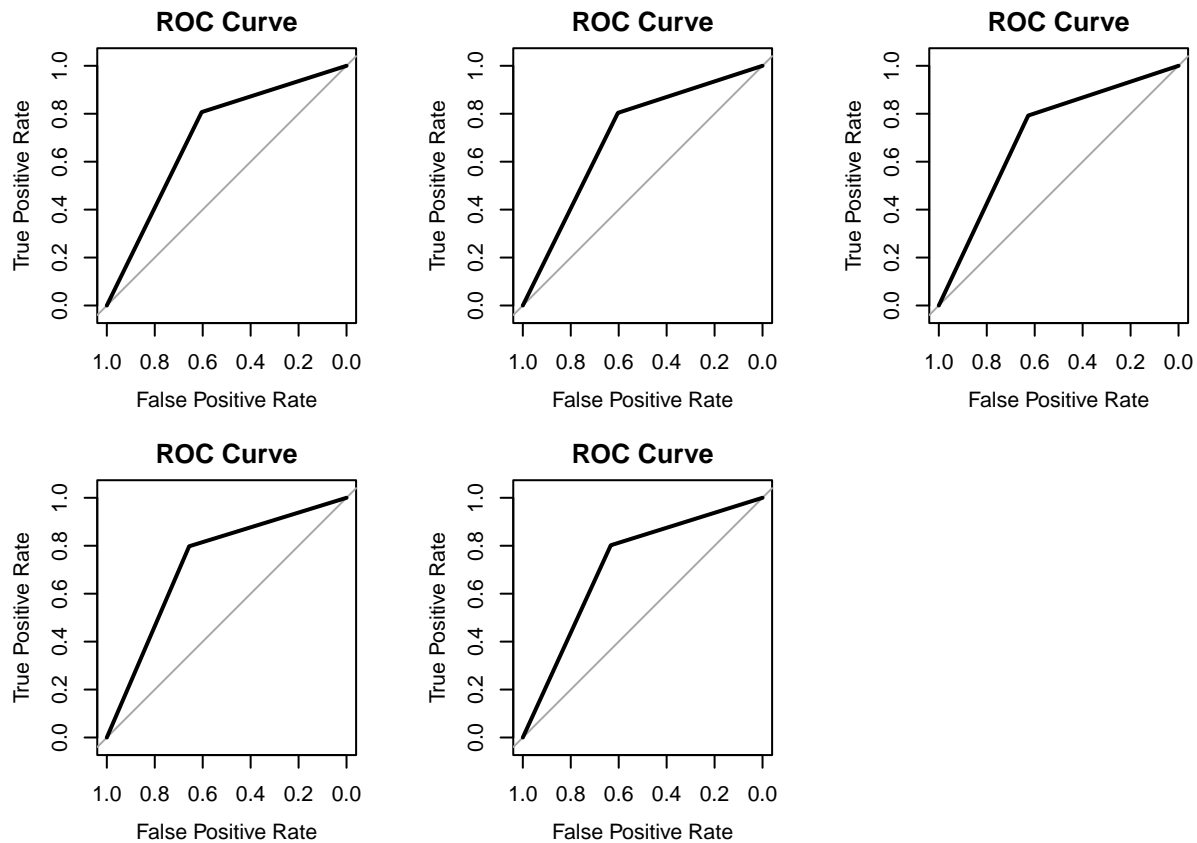
auc(roc(wine$good[d], as.numeric(knn_d)))

## Area under the curve: 0.7271

auc(roc(wine$good[e], as.numeric(knn_e)))

## Area under the curve: 0.7176

```



```
par(mfrow = c(2, 3))

boxplot(wine$volatile.acidity, wine$good, ylab = "Volatile Acidity", xlab = "Wine Quality")


boxplot(wine$residual.sugar, wine$good, ylab = "Residual Sugar", xlab = "Wine Quality")

boxplot(wine$total.sulfur.dioxide, wine$good, ylab = "Total Sulfur Dioxide", xlab = "Wine Quality")


boxplot(wine$free.sulfur.dioxide, wine$good, ylab = "Free Sulfur Dioxide", xlab = "Wine Quality")

boxplot(wine$sulphates, wine$good, ylab = "Sulphates", xlab = "Wine Quality")


boxplot(wine$alcohol, wine$good, ylab = "Alcohol", xlab = "Wine Quality")
```




Wine Quality




Wine Quality




Wine Quality



Wine Quality



Wine Quality



Wine Quality